

Chapter 1

Applications

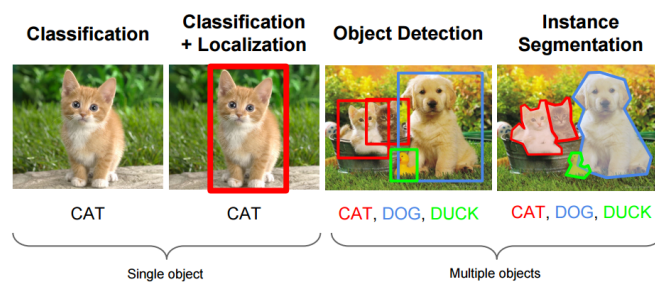


FIGURE 1.1: Some applications examples

1.1 Classification

Train a classification model with softmax loss. The input is the entire image and the output are C probabilities (one per class) of being in the image.




FIGURE 1.2: Classification

1.2 Classification + Localization

Goal is to find a fixed number of objects (one or many) in an image

3 easy steps recipe:




FIGURE 1.3: 1.- Train (or download) a classification model (AlexNet, VGG, GoogLeNet)

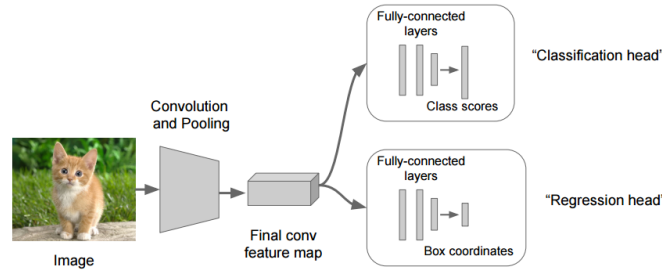


FIGURE 1.4: 2.- Attach a new fully-connected "regression head" to the network to compute bounding boxes (x,y,w,h)

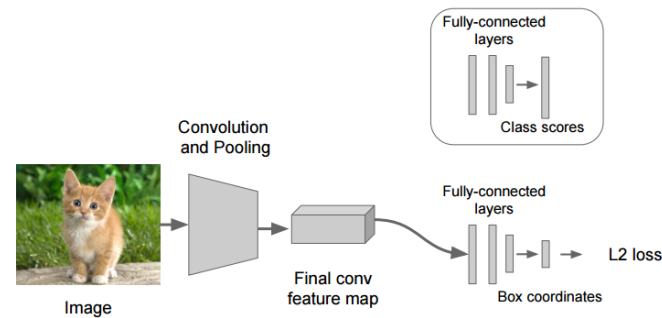


FIGURE 1.5: 3.- Train the regression head only with SGD and L2 loss with the bounding boxes as ground-truth.

When training the regression head there are two options: You can backpropagate only till the regression head or the entire network. 2nd option will improve a little bit the accuracy at a expense of higher training computation cost. If you choose option 2 you will be changing the original Conv layers on which the classification head is trained. So there are two options. Or you have two independent networks: the original one (Conv+classification head) and the other one (Modified Conv+regression head). Or you train both at the same time so you will have only one model(classification head + regression head).

The final DNN for classification + uses both heads at test time

- Classification head:
- The output are C numbers (one per class)
- Regression head:
- There are two regression heads strategies (choose one):

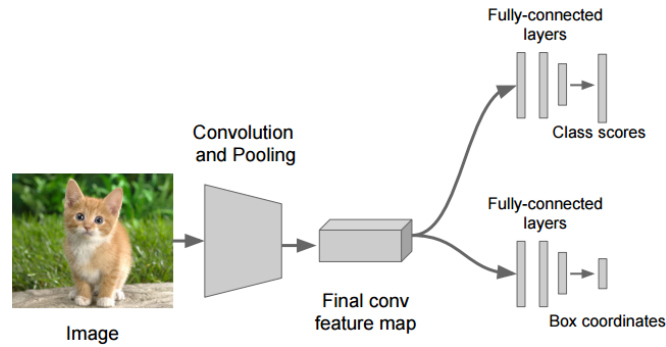


FIGURE 1.6: 4.- At test time use both heads

1. Class agnostic: Output 4 numbers (one bounding box)
2. Class specific: Output $C \times 4$ numbers (one bounding box per class)

Notice that instead of training a regression head which outputs bounding boxes, we could have train it to output a specific number of joints or whatever

Sliding windows (DO NOT DO THIS)

At test time this (classification + regression) network should be applied inside a sliding windows at multiple scales. (This would have an **enormous computation cost**, we will see later that there is a better way of doing this).




FIGURE 1.7: Classification + Localization - Box regression and Sliding windows diagram

The idea is to run the classification + regression network inside a sliding windows and at multiple scales. So for each windows position we obtain a cat bb (x,y,w,h) and the probability of a cat being inside the current sliding windows position. (This would have an enormous computation cost, we will see later that there is a better way of doing this)

Finally, we have obtained multiple bound boxes and we must merge them and estimate the prob of this final bb of containing a cat. To do so, an algorithm such us non-max suppression could be used.

Bounding Box regression (ONLY INPUT IMAGE FIXED SIZE)

Fully connected layers can only deal with input of a fixed size, because it requires a certain amount of parameters to "fully connect" the input and output.




FIGURE 1.8: Classification + Localization - Box regression architecture

In the example network with fully-connected layers at the end, a 221×221 image will output a 1000 size vector of class scores. If we apply the network on a larger image, the network will fail because of the inconsistency between the input and parameters of the first fully-connected layer.

In other words, we would have to execute the DNN in each sliding windows position. This has an enormous computation cost. To solve this we can transform the FC layers to Conv layers. (see right column)

Without sliding windows (DO THIS)

Any FC layer can be converted to a CONV layer. For example, in this model, the FC layer with $K=4096$ that is looking at the input volume of size $5 \times 5 \times 1024$ can be equivalently expressed as a CONV layer with $F = 5, P = 0, S = 1, K = 4096$.

In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be $1 \times 1 \times 4096$ since only a single depth column “fits” across the input volume, giving identical result as the initial FC layer.

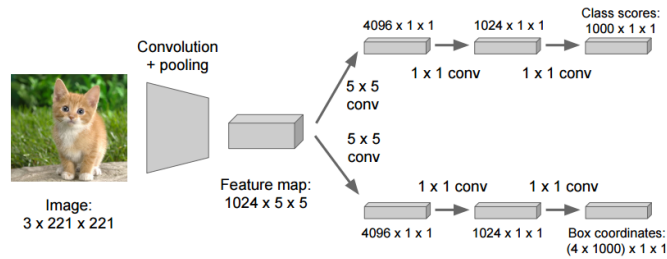


FIGURE 1.9: Classification + Localization - Without sliding windows

FC can only deal with input of a fixed size, while convolutional layers just “slide” the same filters across the input, so it can basically deal with input of an arbitrary spatial size.

At train time we train with samples of the sliding windows. At test time we can evaluate input images of higher dimension sharing computations.




FIGURE 1.10: ConvLayers allow input of any size in an efficient way

For example, in this image, we have trained the model with input image sizes of 14×14 producing and output vector of $1 \times 1 \times C$. At test time, we are using input images of 16×16 which produces and output of $2 \times 2 \times C$. This is equivalent to execute the model 4 times (in the four corners of the image) like in the previous example of the cat. But instead of doing the whole computation 4 times, we are sharing the computation, the extra computation is only done in the yellow parts.

No more sliding windows!!

1.3 Object Detection

Object detection is a harder problem than classification + localization. Now, there can be a non fix number of instances of the same class as well as other classes. We can not use regression with detection (as within classification) because we need variable sized outputs. So we have to go back to sliding windows? It would be very costly to run a classification model in each windows position. The answer is no. We will first search for tiny subsets of possible positions. This models are called R-CNN (Region based CNN).

Evaluation. To evaluate a detection model we use "mean average precision" (mAP). Compute average precision (AP) separately for each class, then average over classes. A detection is a true positive if it has a IoU with a ground-truth box greater than some threshold (usually 0.5) (mAP@0.5). Combine all detection from all test images to draw a precision / recall curve for each class; AP is the area under the curve. mAP is a number from 0 to 100. The higher the better.

There are two main networks used for object detection explained in previous chapters:

- R-CNN (Region Based CNN) - R-CNN, Fast R-CNN & Faster R-CNN
- YOLO (You only look once)

1.4 Segmentation

1.4.1 Semantic Segmentation

Semantic Segmentation tries to label every pixel. It does not differentiate instances (e.g. cows are a massive compact blue mass)

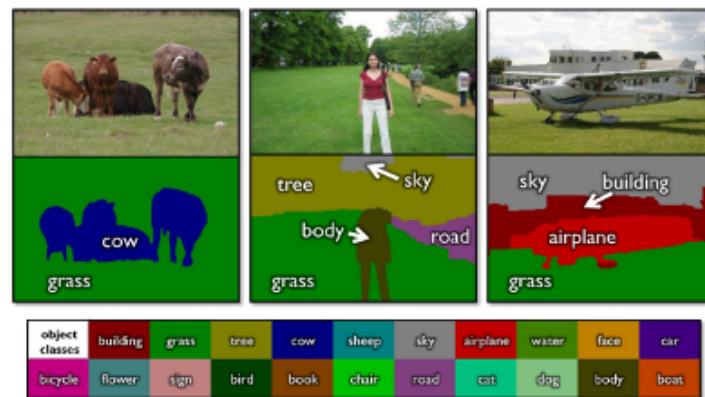


FIGURE 1.11: Semantic Segmentation

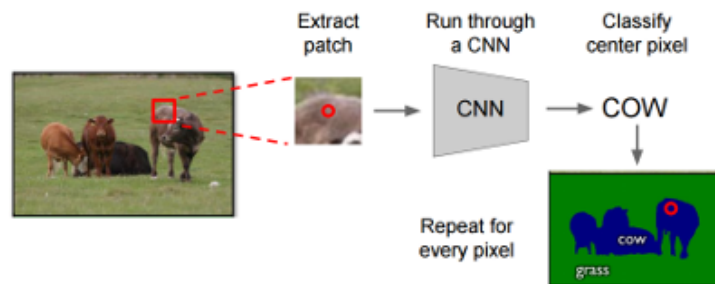


FIGURE 1.12: General idea. Problem: Smaller output due to pooling. Other more complex solutions are required

There are three common strategies used for semantic segmentation:

- Multiscale, or
- Refinement, or
- Upsampling

Multiscale

Solve problem by working in different scales and up-scaling output. In this example they are also doing a parallel process to compute super-pixels to improve the segmentation.




FIGURE 1.13: Farabet et al, “Learning Hierarchical Features for Scene Labeling”, TPAMI 2013

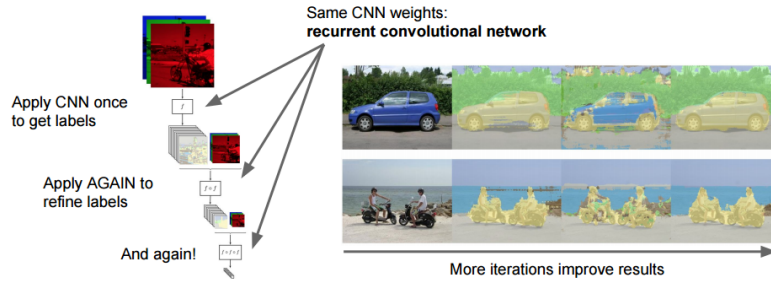


FIGURE 1.14: Pinheiro and Collobert, “Recurrent Convolutional Neural Networks for Scene Labeling”, ICML 2014

Refinement

1. Pass the input image through a CNN which produces the low res image segmentation (labels).
2. Rescale the inputs to the lower resolution and apply again the ConvNet to the lower resolution input image and results from the previous segmentation
3. Goto 2 until threshold

Upsampling

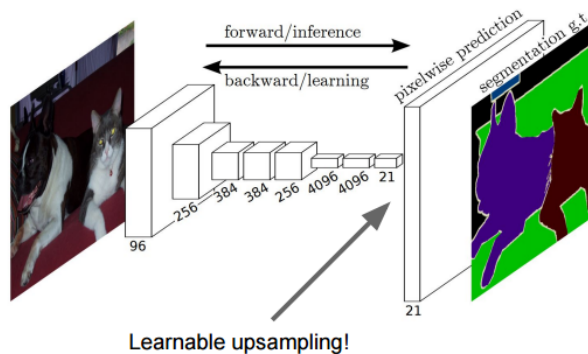


FIGURE 1.15: Long, Shelhamer, and Darrell, “Fully Convolutional Networks for Semantic Segmentation”, CVPR 2015

They pass the input image through a CNN which produces the low res image segmentation (labels). But with this method, they add an extra learnable layer at the end of the ConvNet

that learns to upsample the low resolution segmentation (Upsampling layer). The also do skip connections

1.4.2 Instance Segmentation

Detect instances, give category, label pixels. simultaneous detection and segmentation (SDS). Normally look very similar to the detection models.




FIGURE 1.16: Instance Segmentation

Architecture similar to R-CNN

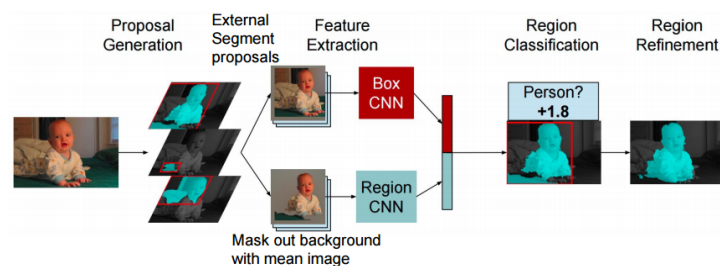


FIGURE 1.17: Architecture - Hariharan et al, “Simultaneous Detection and Segmentation”, ECCV 2014

1. External segment proposals that outputs pixels not boxes
2. Produce a BBox of the segmented region
3. Take the BBox image and run it through a CNN
4. Take the BBox image and set the non segment proposal pixels to the mean image value of the dataset and run it through a CNN
5. Concatenate both features and run it through a classifier
6. Refine the proposed region

Notice that this is very similar to R-CNN

For the refinement step there is a following paper that improves it (Hariharan et al, Hypercolumns for Object Segmentation and Fine-grained Localization, CVPR 2015). The idea is that

once we have the region cropped pass it through AlexNet and extract features of various layers. Then, upsample this feature maps and combine them together. Finally, do a logistic classifier for each of the pixels that predicts how much likely is to be background.

Architecture similar to Faster R-CNN

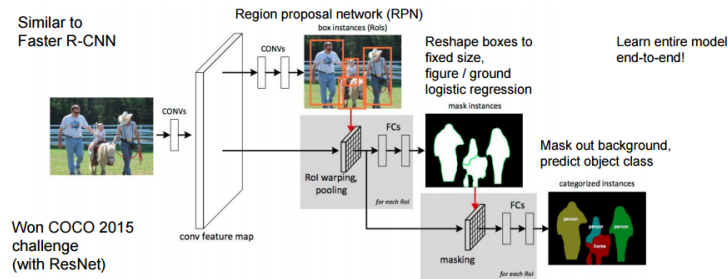


FIGURE 1.18: Architecture - Dai et al, “Instance-aware Semantic Segmentation via Multi-task Network Cascades”, arXiv 2015

Stuck the model in the left to ResNet. Similar to Faster R-CNN

From the high resolution feature map propose regions, then reshape boxes and finally mask background and predict object class.

Notice that the three steps (intermediate levels of the net) can be evaluated with ground-truth.

1.5 Image captioning

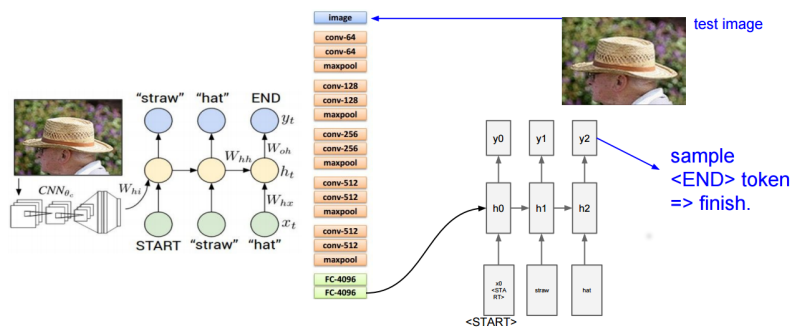


FIGURE 1.19: “Deep Visual-Semantic Alignments for Generating Image Descriptions”, Karpathy and Fei-Fei

Summary

- Localization
 - Find a fixed number of objects (one or many)

- L2 regression from CNN features to box coordinates
- Overfeat: Regression + efficient sliding window with FC \rightarrow conv conversion
- Deeper networks do better
- Object Detection
 - Find a variable number of objects by classifying image regions
 - Before CNNs: dense multiscale sliding window (HoG, DPM)
 - Avoid dense sliding window with region proposals
 - R-CNN: Selective Search + CNN classification / regression
 - Fast R-CNN: Swap order of convolutions and region extraction
 - Faster R-CNN: Compute region proposals within the network
 - Deeper networks do better
- Semantic segmentation
 - Classify all pixels
 - Fully convolutional models, downsample then upsample
 - Learnable upsampling: fractionally strided convolution
 - Skip connections can help
- Instance Segmentation
 - Detect instance, generate mask
 - Similar pipelines to object detection